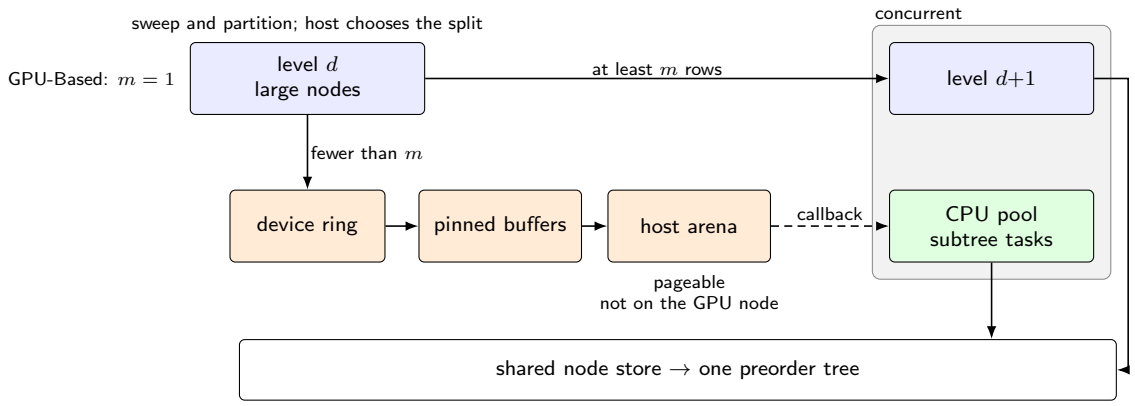
sweep and partition; host chooses the split
GPU-Based: $m = 1$
level d
large nodes
at least m rows
concurrent
level d+1
fewer than m
device ring
pinned buffers
host arena
callback
CPU pool
subtree tasks
pageable
not on the GPU node
shared node store → one preorder tree